

Towards user-friendliness in proof assistants: an effect-based attempt

APRIL GONÇALVES, Metastate AG

Proof assistants provide a framework for modelling and verification of theories, as well as trust-worthy software. However, such power is usually only available to experts. We propose a new approach, based on algebraic effects and handlers, to integrate different automated proof strategies that enables newcomers to take advantage of proof assistants without a more in-depth understanding of underlying theory. Our approach gives beginner users an effect system, a handful of effectful strategies (tactics, proof search and a SMT solver) and their handlers under a shared interface, while advanced users can extend our system with new effects and new handlers. Lastly, we prototype the system as a library in Agda. While our prototype is minimal, it shows how easily proofs can be carried on so long as the user has the correct intuition. We believe our system empowers non-experts and has the potential to bring verified software to many relevant industries, such as finance.

Additional Key Words and Phrases: effect handlers, proof assistants, user experience

ACM Reference Format:

April Gonçalves. 2020. Towards user-friendliness in proof assistants: an effect-based attempt. 1, 1 (September 2020), 4 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 VERIFICATION FOR THE MASSES

The usability of proof assistants such as Coq, Agda, Isabelle/HOL and Twelf is hard to measure: all of them require significant domain knowledge that frequently matches a deeper understanding of Type Theory, the theory that enables such assistants to exist. Coq and Isabelle/HOL are more commonly applied outside their niche, and their programming style is sharply different from mainstream programming and even functional programming. Granted, they come with their own integrated development environment (IDE), which may make it easier for students and newcomers. To our knowledge, no usability study has been held for a proof assistant, however there exists studies that account for an informal “usability” criteria – they appear to measure how fast users can familiarise themselves with the system or IDE.

Folklore within the proof assistants community seems to show that typical users find Coq hard to use due to a large library of tactics and difficult readability of the code after its completion. Many Coq learning materials suggest users to add comments for structural induction cases and inductive hypothesis. As mentioned, no user studies were conducted to attest such hypotheses.

It also has to be granted that proof assistants are relatively new technology and there are no guidelines or even an intuition on how such systems should perform or what features should be available to the users. The closest model we have are proofs by hand, however, following that model for proof assistants has two main pitfalls: (a) proofs by hand are also not widely taught in mathematics classes during school years; and (b) proofs by hand and automated proof writing

Author’s address: April Gonçalves, april@cyberglot.me, Metastate AG, Edinburgh, UK.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2020 Association for Computing Machinery.

XXXX-XXXX/2020/9-ART \$15.00

<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

employ fundamentally different paradigms, where the latter requires axioms and formulations to be encoded explicitly, commonly known as a “pedantic proof style”.

1.1 Empowering users of different levels of expertise

To popularise software verification, we must make proof writing less of a burden to the user. Currently, only specialists can write proofs for their programs, even though the working programmer understands the domain and invariants of their projects – otherwise they would not be able to write any software. *Our hypothesis is that programmers are well-equipped to derive and prove properties of their software, but they lack the mathematical maturity and vocabulary to carry on a formal proof.* We have evidence that students fail to produce well-made proofs due to the lack of mathematical maturity, even though they do understand the subject matter at hand.

In our approach, we propose an (algebraic) effects and handlers view of such proofs, based on prior work developed by the Andromeda proof assistant. Here, our users will program as they would normally and invoke a proof environment as an effect to prove certain properties as they go. Given that the proof environment is *just* an effect, we envision that different proof styles (e.g, Agda-style dependent types, SMT solver, proof search) can be “composed” under a shared interface, a proof object that can manipulate itself while querying different automated proof engines.

The reasons we employ algebraic effects and handlers are manifold:

- (1) as proofs cannot be fully automated, all approaches that try to automate the process (e.g, proof search, SMT solver) may either be non-deterministic or never find a solution. Therefore, the system should be able to handle “impure” computations and errors. Algebraic effects and handlers have well-defined semantics and provide a simple interface for performing effects. With them, we avoid indiscriminate effects that are often error-prone and intricate effects machinery such as monad transformers;
- (2) its semantics accommodate composition of arbitrary effects and the composition of multiple handlers, which means users have the ability to weaken more general strategies into specific ones while maintaining the original untouched;
- (3) it has well-defined semantics, and it is the main feature of many new research languages, such as Koka, Eff, Frank and Effekt, what suggests a greater potential to be investigated further.

1.2 Goals and Contributions

The present project is still under development, and we intend to develop a core calculus and a full compiler-level implementation into the Juvix¹ programming language.

Goals. The project exists to enable users to write safe, verified code without investing time into learning dependent types and theorem proving. The project’s goals are presented as follows:

- (1) make software verification more accessible to programmers familiar with functional programming;
- (2) foster knowledge sharing between domain experts and regular programmers;
- (3) lower the barrier to entrance into software verification to non-experts.

Contributions. While the project is a work-in-progress, this paper provides three important contributions that lays out the full extent of the work, which are presented as follows:

- (1) a specification of a user interface for automated proofs based on algebraic effects and handlers that permits user-defined effects for proof strategies, and also user-defined handlers for existing effects (Section 2);

¹Juvix is a dependently-typed programming language for smart contracts. More information at juvix.org

- (2) a minimal prototype in Agda that validates the feasibility of such system (Section 3);
- (3) a short technical account of the effort required to materialise such a user interface into proof assistants (Section 3.1).

2 INTRODUCING WITCH

```

data N : Type where
  0      : N
  suc    : N → N

_+_      : N → N → N
0 + n    = n
(suc m) + n = suc (m + n)

+comm    : ∀ (x y : N) → ⊢ x + y ≡ y + x
+comm 0 y = solve { 0 + y ≡ y + 0 } via SMT
+comm (suc x) y = search { x + (suc y) ≡ suc (x + y) }
                  via Backtracking

```

regular Agda-style datatype and function definitions

Witch's effect wrapper

effect operation

effect handler

Fig. 1. Low-fidelity prototype of Witch

Downsides of our approach. There are two main styles for writing proofs within a proof assistant: external verification and internal verification. The latter concerns itself with code that is *correct-by-construction*, meaning that only valid states allowed and it carries all information necessary to make proofs trivial; while the former concerns itself with proofs about non-proof-carrying code. Witch's primarily focus is on external verification. We believe that newcomers will struggle to write proof-aware code, since it departs strongly from mainstream programming, even within the functional style. We expect users to write their code as they normally would do and then prove necessary properties. However, Witch does not support writing correct-by-construction code.

3 WITCH IN AGDA

Why Agda? Agda provides high-level meta-programming constructs, and its coding style is a reminiscent of ML family of functional programming languages.

3.1 Technical Effort Required

Although Witch is implementable at the library level, it requires that proof assistants to provide a set of off-the-shelf features. The most important feature is type-level meta-programming. Witch is

only possible given the proof assistant's ability to manipulate its own proof terms to generate new variables, holes, goals, among others and manipulate terms and types in a proof-carrying fashion. Secondly, the proof assistant should feature an interface with the outside world. Although it is not inherently a requirement, its absence means that no solver will be integrated, diminishing the proposed value of the tool. Lastly, proof assistants must support algebraic effects and handlers. If there is no built-in solution, a library-level solution is possible – as we did in Agda. However, a solid implementation may not be trivial to achieve. In Agda, we translate all calls into optimised continuation-passing style that inlines all calls. This strategy also works at compiler-level if necessary. There are many strategies to efficiently implement algebraic effects and handlers, but we went for the one that seemed the most straight-forward method to implement it in Agda.